


Deep Learning: Real World Applications and Implementation Details

Webinar Series on Applied Artificial Intelligence
Vikram Sarabhai Space Centre



Litu Rout
Space Applications Centre
Indian Space Research Organisation



Three Pillars of Deep Learning



Three Pillars of Deep Learning

- Setting Up DL Environment
- Defining Problem Statement
- Implementation Details

Setting Up DL Environment



- Data Processing
- Network Design
- Visualization

Setting Up DL Environment

- Data Processing
- Network Design
- Visualization

Data Processing


- Programming Language
- Image Processing
- Numeric Computation
- Data Manipulation



Data Processing



- Programming Language
- Image Processing
- Numeric Computation
- Data Manipulation



Data Processing

- Package Installation via “pip”
>> pip install package
- Package Installation via “conda”
>> conda install package



Many packages ship pre-installed in Anaconda

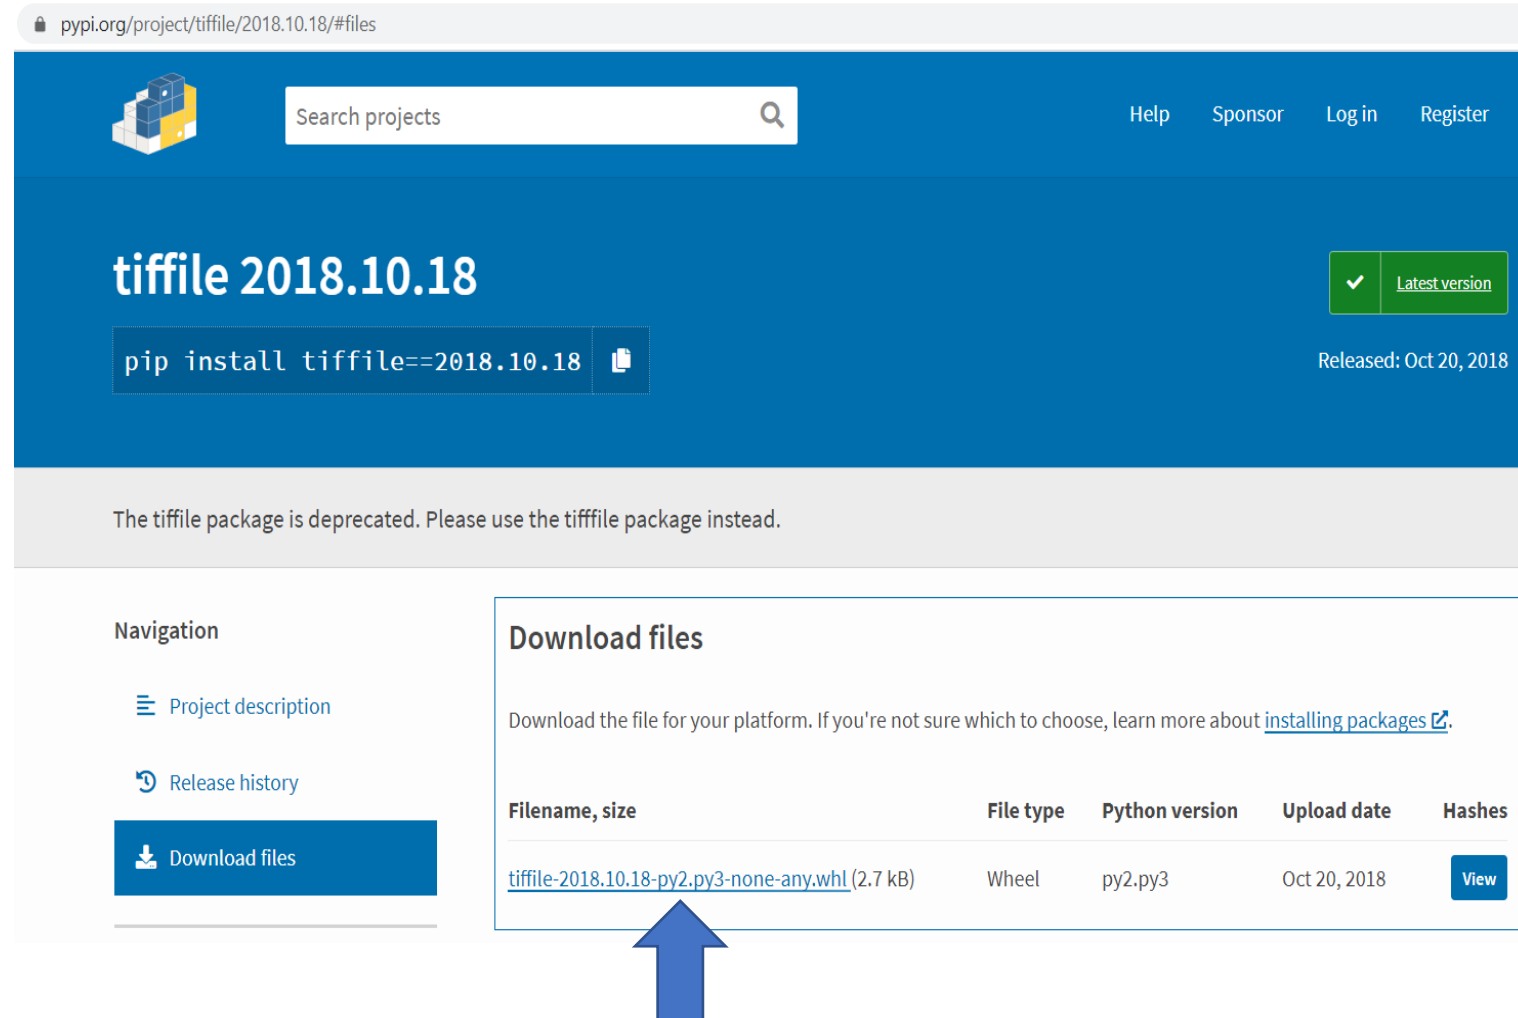
Data Processing

- Offline Installation
 - Download on Thin Client

>> pip install package.whl

or

>> pip install package.tar.gz



The screenshot shows the PyPI project page for **tiffle 2018.10.18**. The page includes a search bar, navigation links (Help, Sponsor, Log in, Register), and a green button labeled "Latest version". Below the package name, there is a code block showing the command `pip install tiffle==2018.10.18`. A message states: "The tiffle package is deprecated. Please use the tiffle package instead." The "Download files" section contains a table with columns: Filename, size, File type, Python version, Upload date, and Hashes. The table lists the file [tiffle-2018.10.18-py2.py3-none-any.whl](#) (2.7 kB) as a Wheel for Python version py2.py3, uploaded on Oct 20, 2018. A blue arrow points to this download link.

Filename, size	File type	Python version	Upload date	Hashes
tiffle-2018.10.18-py2.py3-none-any.whl (2.7 kB)	Wheel	py2.py3	Oct 20, 2018	View

Setting Up DL Environment


- Data Processing
- **Network Design**
- Visualization

Network Design

- Popular Libraries



Caffe




PYTORCH




dy/net



PyTorch (Solid) vs TensorFlow (Dotted) Raw Counts



DL Libraries in 2018



Setting Up DL Environment

- Data Processing
- Network Design
- Visualization


Visualization

- Popular Libraries


matplotlib

Seaborn

plotly



Visualization of Loss Surface



Jin et al. 2020, ICML

Three Pillars of Deep Learning

- Setting Up DL Environment
- **Defining Problem Statement**
- Implementation Details

Defining Problem Statement

- Linear Function Approximation

- Dataset Preparation

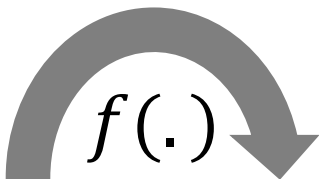
$$\{(x_p, y_p)\}_{p=1}^n \subset R^{d_{in} \times d_{out}}$$

- Function Approximator

$$f(m, c, x) = m x + c$$

- Goal

$$m = ?, c = ?$$



x	y
-10	-48
-9	-43
...	...
10	52

Defining Problem Statement

- Linear Function Approximation


- Error Computation

$$l(f(m, c, x), y) = \frac{1}{2} (f(m, c, x) - y)^2$$

$$L(m, c) = \frac{1}{2n} \sum_{p=1}^n (f(m, c, x_p) - y_p)^2$$

- Optimization

$$(m^*, c^*) = \arg \min_{(m, c) \in \mathbb{R}^{1 \times 1}} L(m, c)$$



Defining Problem Statement

- Linear Function Approximation

- Optimization

$$(m^*, c^*) = \arg \min_{(m, c) \in \mathbb{R}^{1 \times 1}} L(m, c)$$

- Learning Algorithm: GD

$$\frac{dm}{dt} = -\eta \frac{\partial L(m, c)}{\partial m(t)}, \quad m(t+1) = m(t) - \eta \frac{\partial L(m, c)}{\partial m(t)}$$

$$\frac{dc}{dt} = -\eta \frac{\partial L(m, c)}{\partial c(t)}, \quad c(t+1) = c(t) - \eta \frac{\partial L(m, c)}{\partial c(t)}$$

Defining Problem Statement


- Linear Function Approximation

- Gradient Descent (GD)

$$m(t + 1) = m(t) - \eta \frac{\partial L(m, c)}{\partial m(t)}$$



$$\frac{\partial L(m, c)}{\partial m(t)} < 0 \quad \text{at point A}$$

$$\frac{\partial L(m, c)}{\partial m(t)} > 0 \quad \text{at point B}$$





Defining Problem Statement

- Linear Function Approximation
 - Gradient Descent (GD)



Defining Problem Statement

- Linear Function Approximation
 - Gradient Descent (GD)



Defining Problem Statement

- Linear Function Approximation

- Gradient Descent (GD)


$$\eta = \beta^{-1}$$

$$m_{t+1} = m_t - \beta^{-1} \nabla L(m_t), \nabla L(m_t) \triangleq \frac{\partial L(m, c)}{\partial m(t)}$$

- Assumption: β -Smoothness

$$\|\nabla L(m_{t+1}) - \nabla L(m_t)\|_2 \leq \beta \|m_{t+1} - m_t\|_2$$

$$\text{Or} \quad \|\nabla^2 L(m_t)\|_2 \leq \beta$$



Defining Problem Statement


- Linear Function Approximation

- Gradient Descent (GD)

$$m_{t+1} = m_t - \beta^{-1} \nabla L(m_t)$$

- Taylor Series Expansion

$$L(m_{t+1}) = L(m_t) + \langle \nabla L(m_t), m_{t+1} - m_t \rangle + \frac{1}{2} (m_{t+1} - m_t)^T \nabla^2 L(m_t) (m_{t+1} - m_t)$$



Defining Problem Statement


- Linear Function Approximation

- Gradient Descent (GD)

$$m_{t+1} = m_t - \beta^{-1} \nabla L(m_t)$$

- Taylor Series Expansion

$$\begin{aligned} L(m_{t+1}) &= L(m_t) + \langle \nabla L(m_t), m_{t+1} - m_t \rangle + \frac{1}{2} (m_{t+1} - m_t)^T \nabla^2 L(m_t) (m_{t+1} - m_t) \\ &\leq L(m_t) + \langle \nabla L(m_t), m_{t+1} - m_t \rangle + \frac{\beta}{2} \|m_{t+1} - m_t\|_2^2, \because \|\nabla^2 L(m_t)\|_2 \leq \beta \end{aligned}$$



Defining Problem Statement


- Linear Function Approximation

- Gradient Descent (GD)

$$m_{t+1} = m_t - \beta^{-1} \nabla L(m_t)$$

- Taylor Series Expansion

$$\begin{aligned} L(m_{t+1}) &= L(m_t) + \langle \nabla L(m_t), m_{t+1} - m_t \rangle + \frac{1}{2} (m_{t+1} - m_t)^T \nabla^2 L(m_t) (m_{t+1} - m_t) \\ &\leq L(m_t) + \langle \nabla L(m_t), m_{t+1} - m_t \rangle + \frac{\beta}{2} \|m_{t+1} - m_t\|_2^2, \because \|\nabla^2 L(m_t)\|_2 \leq \beta \\ &\leq L(m_t) - \beta^{-1} \|\nabla L(m_t)\|_2^2 + \frac{\beta^{-1}}{2} \|\nabla L(m_t)\|_2^2, \end{aligned}$$



Defining Problem Statement


- Linear Function Approximation

- Gradient Descent (GD)

$$m_{t+1} = m_t - \beta^{-1} \nabla L(m_t)$$

- Taylor Series Expansion

$$L(m_{t+1}) \leq L(m_t) - \frac{\beta^{-1}}{2} \|\nabla L(m_t)\|_2^2$$



Defining Problem Statement

- Linear Function Approximation


- Iteration Complexity

$$L(m_1) \leq L(m_0) - \frac{\beta^{-1}}{2} \|\nabla L(m_0)\|_2^2$$

$$L(m_2) \leq L(m_1) - \frac{\beta^{-1}}{2} \|\nabla L(m_1)\|_2^2$$

⋮

$$L(m_T) \leq L(m_{T-1}) - \frac{\beta^{-1}}{2} \|\nabla L(m_{T-1})\|_2^2$$



Defining Problem Statement

- Linear Function Approximation


- Iteration Complexity

$$\cancel{L(m_1)} \leq L(m_0) - \frac{\beta^{-1}}{2} \|\nabla L(m_0)\|_2^2$$

$$\cancel{L(m_2)} \leq \cancel{L(m_1)} - \frac{\beta^{-1}}{2} \|\nabla L(m_1)\|_2^2$$

⋮

$$L(m_T) \leq \cancel{L(m_{T-1})} - \frac{\beta^{-1}}{2} \|\nabla L(m_{T-1})\|_2^2$$




Defining Problem Statement

- Linear Function Approximation
 - Iteration Complexity

$$L(m_T) \leq L(m_0) - \frac{\beta^{-1}}{2} \sum_{t=0}^{T-1} \|\nabla L(m_t)\|_2^2$$

$$L(m_0) - L(m_T) \geq \frac{\beta^{-1}}{2} \sum_{t=0}^{T-1} \|\nabla L(m_t)\|_2^2$$




Defining Problem Statement

- Linear Function Approximation

- ϵ -Stationary Solution

$\|\nabla L(m_T)\|_2 \leq \epsilon \implies$ For all $t = 0, \dots, T - 1$,

$$\|\nabla L(m_t)\|_2 > \epsilon$$



Defining Problem Statement

- Linear Function Approximation


- ϵ -Stationary Solution

$\|\nabla L(m_T)\|_2 \leq \epsilon \implies$ For all $t = 0, \dots, T - 1$,

$$\|\nabla L(m_t)\|_2 > \epsilon$$

- Iteration Complexity

$$L(m_0) - L(m_T) \geq \frac{\beta^{-1}}{2} \sum_{t=0}^{T-1} \|\nabla L(m_t)\|_2^2$$



Defining Problem Statement

- Linear Function Approximation


- ϵ -Stationary Solution

$\|\nabla L(m_T)\|_2 \leq \epsilon \implies$ For all $t = 0, \dots, T - 1$,

$$\|\nabla L(m_t)\|_2 > \epsilon$$

- Iteration Complexity

$$L(m_0) - L(m_T) \geq \frac{\beta^{-1}}{2} \sum_{t=0}^{T-1} \epsilon^2$$



Defining Problem Statement

- Linear Function Approximation


- ϵ -Stationary Solution

$\|\nabla L(m_T)\|_2 \leq \epsilon \implies$ For all $t = 0, \dots, T - 1$,

$$\|\nabla L(m_t)\|_2 > \epsilon$$

- Iteration Complexity

$$L(m_0) - L(m_T) \geq \frac{\beta^{-1}}{2} T \epsilon^2$$



Defining Problem Statement

- Linear Function Approximation

- ϵ -Stationary Solution


$$\|\nabla L(m_T)\|_2 \leq \epsilon \implies \text{For all } t = 0, \dots, T - 1,$$

$$\|\nabla L(m_t)\|_2 > \epsilon$$

- Iteration Complexity

$$L(m_0) - L(m_T) \geq \frac{\beta^{-1}}{2} T \epsilon^2$$

$$T \leq \frac{2\beta(L(m_0) - L(m_T))}{\epsilon^2} = \mathcal{O}\left(\frac{1}{\epsilon^2}\right)$$




Defining Problem Statement

- Linear Function Approximation
 - Stochastic Gradient Descent (SGD)

$$L(m, c) = \frac{1}{2n} \sum_{p=1}^n (f(m, c, x_p) - y_p)^2$$

$$L_{\mathcal{B}}(m, c) = \frac{1}{2|\mathcal{B}|} \sum_{p=1}^{|\mathcal{B}|} (f(m, c, x_p) - y_p)^2$$

$$m(t + 1) = m(t) - \eta \frac{\partial L_{\mathcal{B}}(m, c)}{\partial m(t)}$$



Defining Problem Statement

- Linear Function Approximation

$$l_p(m, c) = (f(m, c, x_p) - y_p)^2$$

- Stochastic Gradient Descent (SGD)

$$L(m, c) = \frac{1}{2n} \sum_{p=1}^n (f(m, c, x_p) - y_p)^2$$

$$L(m, c) = \frac{1}{2n} (l_1(m, c) + l_2(m, c) + \cdots + l_n(m, c))$$

Defining Problem Statement

- Linear Function Approximation

$$l_p(m, c) = (f(m, c, x_p) - y_p)^2$$

- Stochastic Gradient Descent (SGD)

$$L(m, c) = \frac{1}{2n} \sum_{p=1}^n (f(m, c, x_p) - y_p)^2$$

$$L(m, c) = \frac{1}{2n} (l_1(m, c) + l_2(m, c) + \cdots + l_n(m, c))$$

$$L_{\mathcal{r}}(m, c) = \frac{1}{2} (l_{\mathcal{r}}(m, c))$$

Defining Problem Statement

- Linear Function Approximation

$$l_p(m, c) = (f(m, c, x_p) - y_p)^2$$

- Stochastic Gradient Descent (SGD)

$$L(m, c) = \frac{1}{2n} \sum_{p=1}^n (f(m, c, x_p) - y_p)^2$$

$$L(m, c) = \frac{1}{2n} (l_1(m, c) + l_2(m, c) + \cdots + l_n(m, c))$$

$$L_{\mathcal{r}}(m, c) = \frac{1}{2} (l_{\mathcal{r}}(m, c))$$

$$m(t + 1) = m(t) - \eta \frac{\partial L_{\mathcal{r}}(m, c)}{\partial m(t)}$$

Defining Problem Statement

- Linear Function Approximation $l_p(m, c) = (f(m, c, x_p) - y_p)^2$
 - Stochastic Gradient Descent (SGD)

$$L(m, c) = \frac{1}{2n} (l_1(m, c) + l_2(m, c) + \cdots + l_n(m, c))$$

$$L_{\mathcal{B}}(m, c) = \frac{1}{2|\mathcal{B}|} \sum_{p=1}^{|\mathcal{B}|} l_p(m, c)$$

$$m(t + 1) = m(t) - \eta \frac{\partial L_{\mathcal{B}}(m, c)}{\partial m(t)}$$

Three Pillars of Deep Learning


- Setting Up DL Environment
- Defining Problem Statement
- **Implementation Details**

Implementation Details

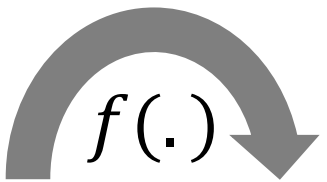
- Linear Function Approximator
- One Layer Neural Network Function Approximator
- Two Layer Neural Network Function Approximator
- Three Layer Convolutional Neural Network Function Approximator

Implementation Details

- Linear Function Approximation





Paired Training Data



x	y
-10	-48
-9	-43
...	...
10	52


Implementation Details

- Linear Function Approximation




Implementation Details

- Linear Function Approximation



Inference Stage



Implementation Details

- Linear Function Approximator
- **One Layer Neural Network Function Approximator**
- Two Layer Neural Network Function Approximator
- Three Layer Convolutional Neural Network Function Approximator

Implementation Details

- One Layer Neural Network Function Approximator


- Dataset Preparation

$$\{(x_p, y_p)\}_{p=1}^n \subset R^{d_{in} \times d_{out}}$$

$$d_{in} = 28 \times 28 = 784$$


$$d_{out} = 10$$

$$n = 60000$$



Implementation Details

- One Layer Neural Network Function Approximator



$f(.)$	
$x = (x_1, x_2, \dots, x_{784})$	$y = (y_1, y_2, \dots, y_{10})$
$(0, 0.5, \dots, 1)$	$(1, 0, \dots, 0)$
$(0.8, 1, \dots, 0)$	$(0, 1, \dots, 0)$
...	...
$(1, 0, \dots, 0.2)$	$(0, 0, \dots, 1)$

Implementation Details


- One Layer Neural Network Function Approximator
 - Function Approximator

$$f_i(\mathbf{m}, c, \mathbf{x}) = m_1 x_1 + m_2 x_2 + \cdots + m_{784} x_{784} + c_i$$

$$= \sum_{j=1}^{784} m_j x_j + c_j$$

$$\mathbf{f}(f_1, f_2, \dots, f_{10}) = \mathbf{M}\mathbf{x} + \mathbf{c}$$

$$[10 \times 1] = [10 \times 784][784 \times 1] + [10 \times 1]$$



Implementation Details

- One Layer Neural Network Function Approximator
 - Function Approximator


$$f_i(\mathbf{m}, c, \mathbf{x}) = m_1x_1 + m_2x_2 + \cdots + m_{784}x_{784} + c_i$$

Trainable Parameters
10x784+10 = 7850

$$= \sum_{j=1}^{784} m_j x_j + c_j$$

$$\mathbf{f}(\mathbf{f}_1, \mathbf{f}_2, \dots, \mathbf{f}_{10}) = \mathbf{M}\mathbf{x} + \mathbf{c}$$

$$[10 \times 1] = [10 \times 784][784 \times 1] + [10 \times 1]$$




Implementation Details

- Linear Function Approximator
- One Layer Neural Network Function Approximator
- **Two Layer Neural Network Function Approximator**
- Three Layer Convolutional Neural Network Function Approximator

Implementation Details

- Two Layer Neural Network Function Approximator
 - Function Approximator


$$f(f_1, f_2, \dots, f_{10}) = M_2(M_1 \mathbf{x} + c_1) + c_2$$




Implementation Details

- Two Layer Neural Network Function Approximator
 - Function Approximator

$$f(f_1, f_2, \dots, f_{10}) = M_2(\sigma(M_1 \mathbf{x} + c_1)) + c_2$$



$$\sigma(x) = x, \text{ if } x \geq 0 \\ = 0, \text{ otherwise}$$





Implementation Details

- Linear Function Approximator
- One Layer Neural Network Function Approximator
- Two Layer Neural Network Function Approximator
- Three Layer Convolutional Neural Network Function Approximator

Implementation Details

- Three Layer Convolutional Neural Network Function Approximator

- Local Connectivity
- Weight Sharing



Implementation Details

- Linear Function Approximator
- One Layer Neural Network Function Approximator
- Two Layer Neural Network Function Approximator
- Three Layer Convolutional Neural Network Function Approximator

Real World Applications

S2A: Wasserstein GAN with Spatio-Spectral Laplacian Attention for Multi-Spectral Band Synthesis

CVPR-EarthVision 2020

Litu Rout


Joint work with Indranil Misra, S Manthira Moorthi and Debajyoti Dhar

Super-resolution as conditional band synthesis

LR-SWIR



HR-SWIR



- Direct super-resolution is intractable.
- Lack necessary geometric attributes.



FCC: NIR (R), R (G), G(B)


- Reformulate as conditional band synthesis.
- Geometry from existing high resolution bands: HR-NIR, R, G.
- Radiometry from corresponding low resolution band: LR-SWIR.

Super-resolution as conditional band synthesis


LR-SWIR



HR-SWIR



- Direct super-resolution is intractable.
- Lack necessary geometric attributes.



FCC: NIR (R), R (G), G(B)


- Reformulate as conditional band synthesis.
- Geometry from existing high resolution bands: HR-NIR, R, G.
- Radiometry from corresponding low resolution band: LR-SWIR.

Traditional Approach


LR-SWIR




Traditional Approach




Traditional Approach




Traditional Approach



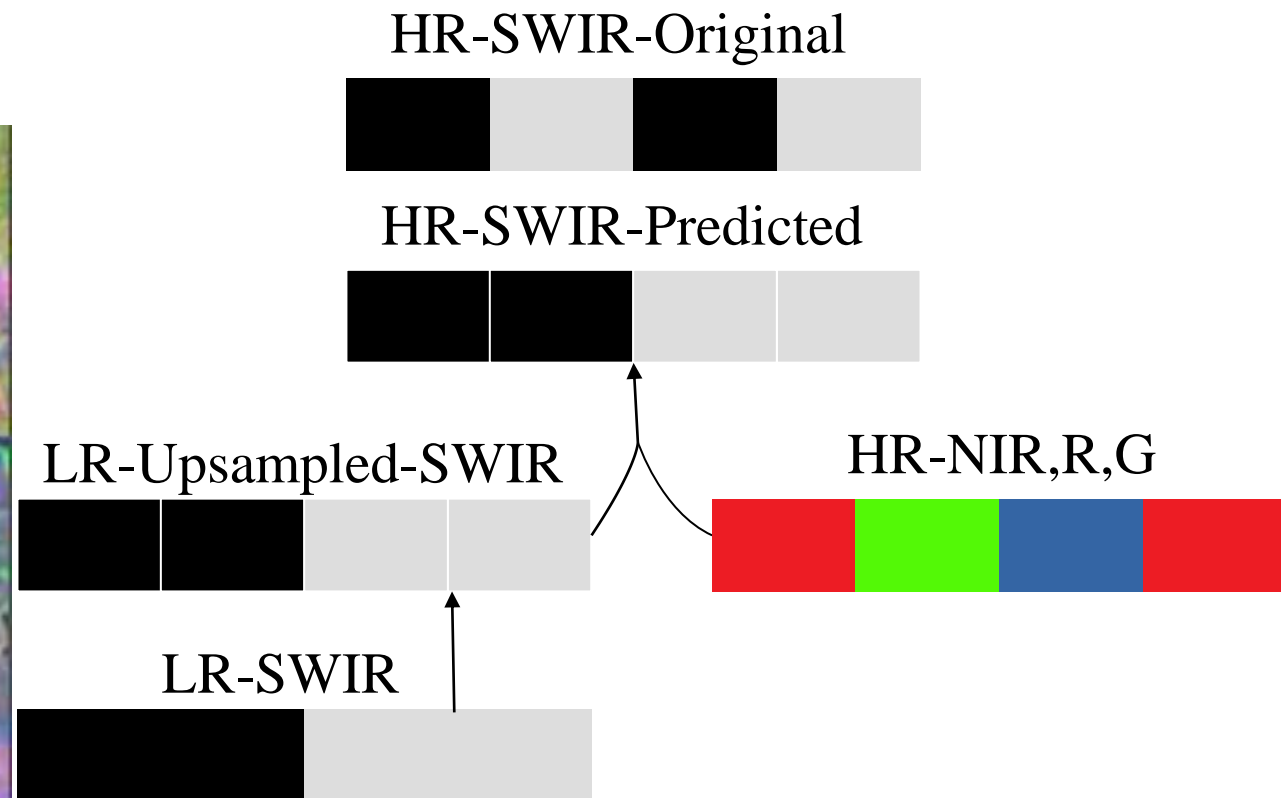
Traditional Approach



Traditional Approach




FCC: SWIR (R), NIR (G), Red (B)



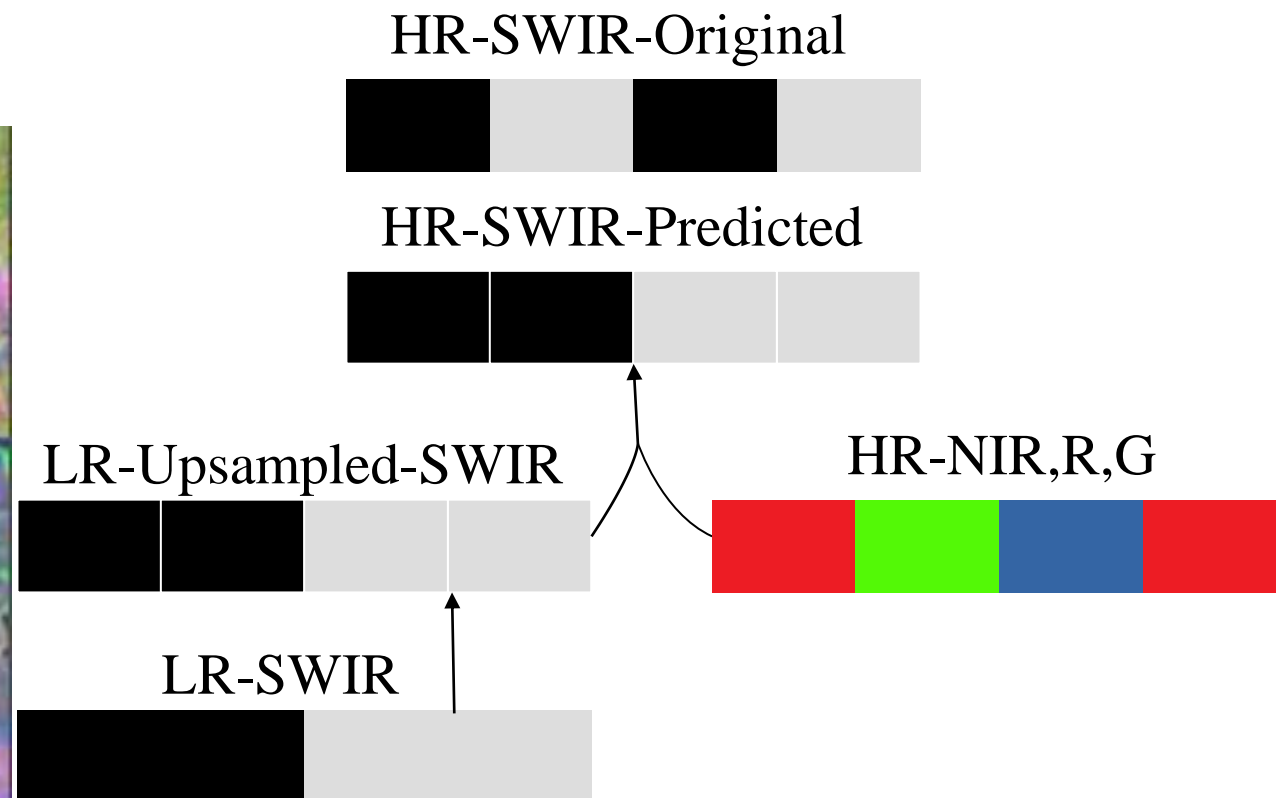
Over dependence on upsampled coarse resolution band results in unpleasant artifacts.

- Geometric distortion
- Radiometric imbalance

Traditional Approach




FCC: SWIR (R), NIR (G), Red (B)



Over dependency on upsampled coarse resolution band results in unpleasant artifacts.

- Geometric distortion
- Radiometric imbalance


Traditional Approach



Over dependency on upsampled coarse resolution band results in unpleasant artifacts.


- Geometric distortion
- Radiometric imbalance

Proposed Approach




Over dependency on upsampled coarse resolution band can be suppressed by replacing it with spatial attention map.

Traditional Approach

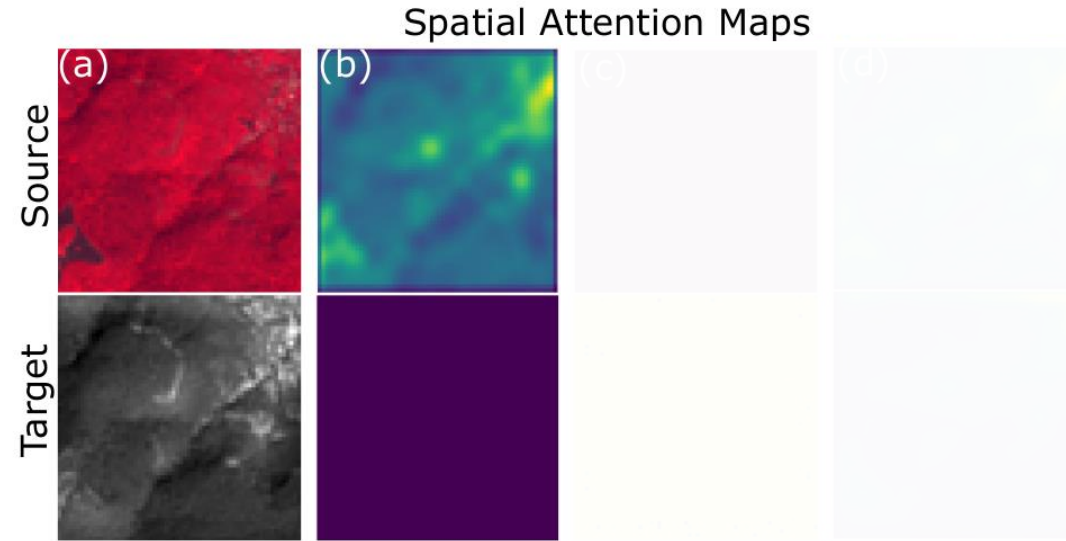
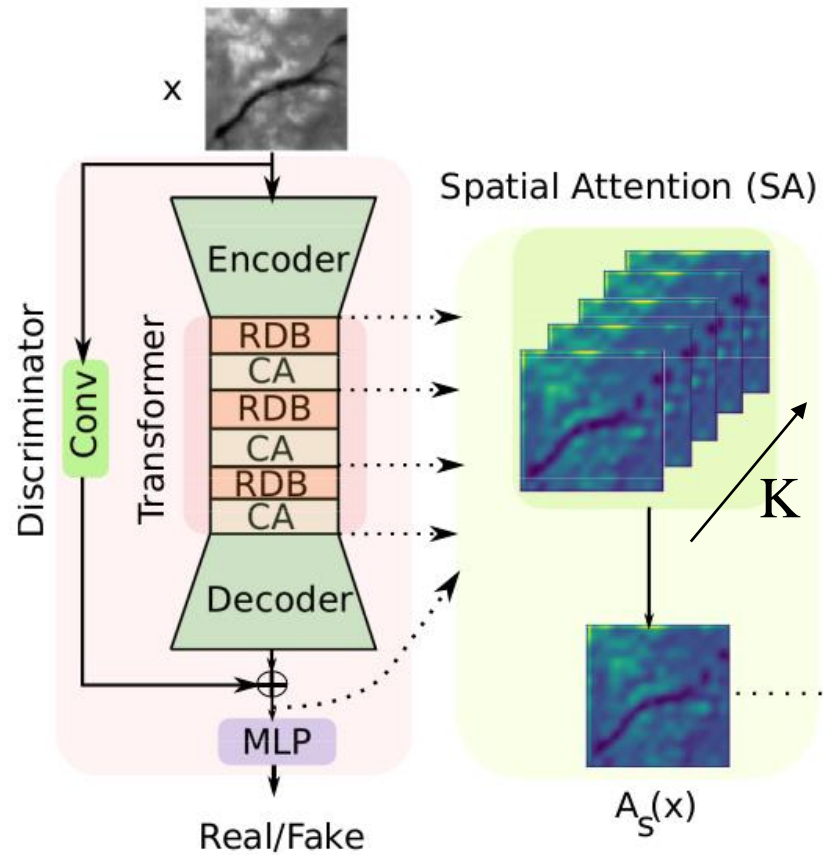


Proposed Approach

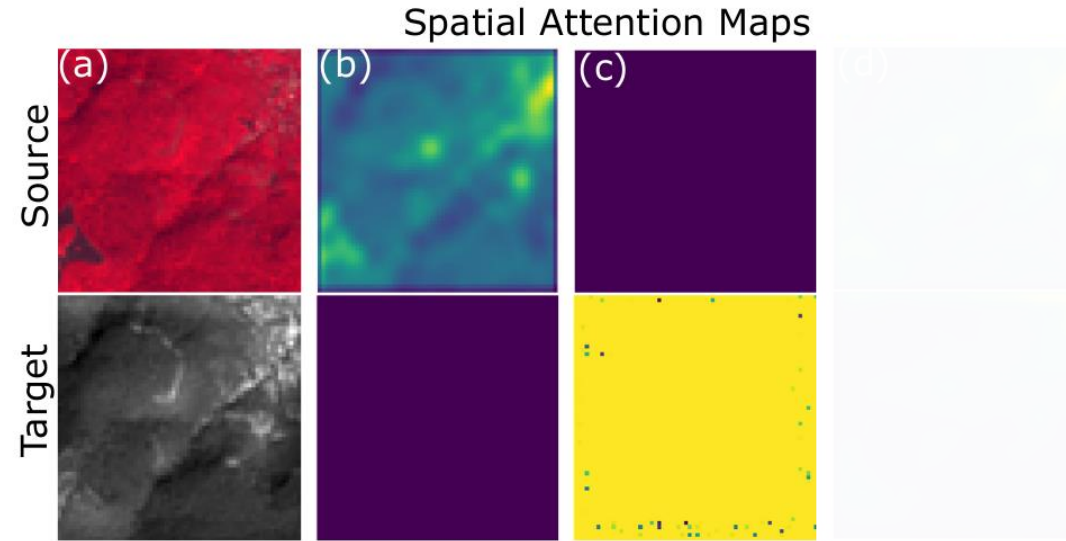



FCC: SWIR (R), NIR (G), Red (B)

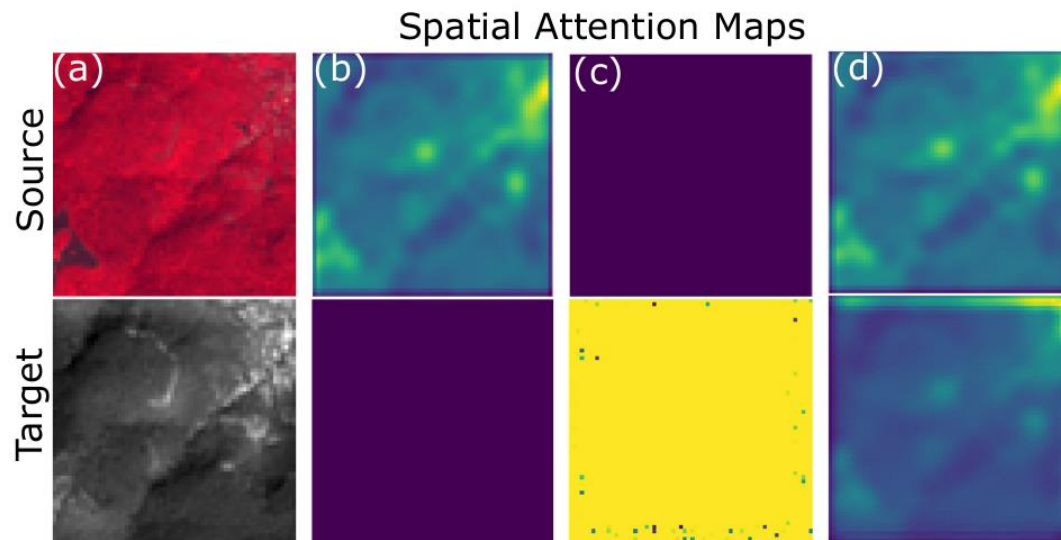

Spatial Attention from Discriminator



Spatial Attention from Discriminator



Spatial Attention from Discriminator




$$A_S(x) = \mathcal{N}(D_S(x)),$$

$$D_S(x) = \sum_{i=1}^K \mathcal{N} \left(\sum_{j=1}^C |A_{ij}(x)| \right)$$

Spatial Attention from Discriminator

Spatial Attention Loss

$$\mathcal{L}_{sa} = \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}, y \sim \mathbb{P}_y} \left[\|A_s(\hat{x}) - A_s(y)\|_2^2 \right]$$



Domain Adaptation Loss

$$\mathcal{L}_{da} = \mathbb{E}_{\tilde{y} \sim \mathbb{P}_{\tilde{y}}, y \sim \mathbb{P}_y} \left[\|A_s(\tilde{y}) - A_s(y)\|_2^2 \right]$$

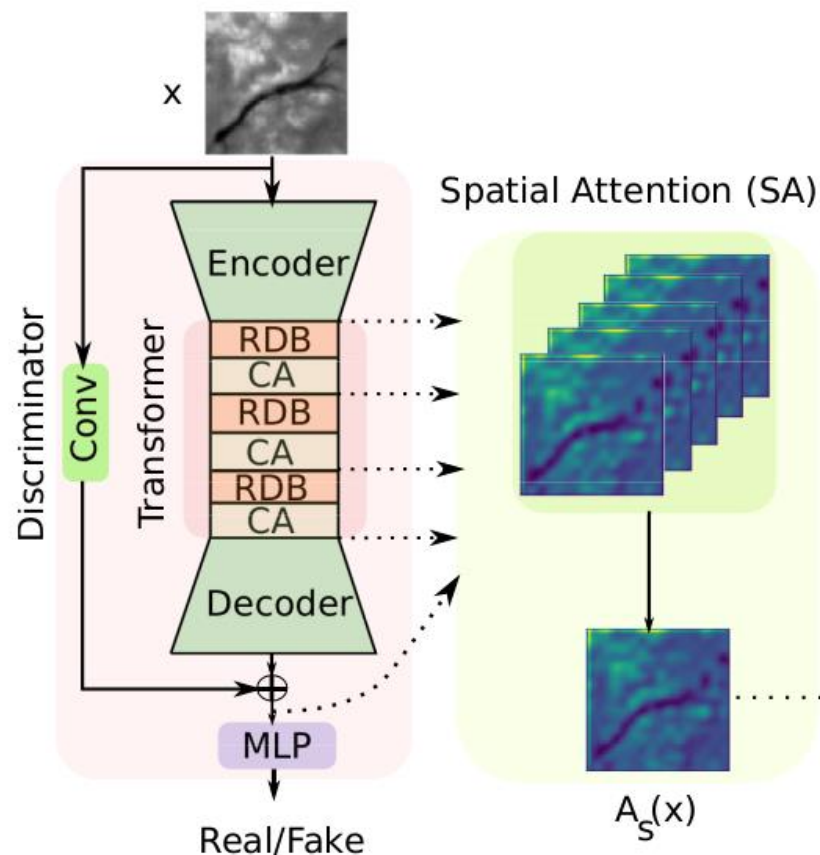
Spatial Attention from Discriminator

Spatial Attention Loss

$$\mathcal{L}_{sa} = \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}, y \sim \mathbb{P}_y} \left[\|A_s(\hat{x}) - A_s(y)\|_2^2 \right]$$

Domain Adaptation Loss

$$\mathcal{L}_{da} = \mathbb{E}_{\tilde{y} \sim \mathbb{P}_{\tilde{y}}, y \sim \mathbb{P}_y} \left[\|A_s(\tilde{y}) - A_s(y)\|_2^2 \right]$$



Discriminator Objective

$$\begin{aligned} \min_D \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [D(\hat{x})] - \mathbb{E}_{x \sim \mathbb{P}_x} [D(x)] \\ + \lambda_{gp} \mathbb{E}_{\tilde{x} \sim \mathbb{P}_{\tilde{x}}} \left[(\|\nabla_{\tilde{x}} D(\tilde{x})\|_2 - 1)^2 \right] \\ + \lambda_{sa} \mathcal{L}_{sa} + \lambda_{da} \mathcal{L}_{da}, \end{aligned}$$

Spatial Attention from Discriminator

Spatial Attention Loss

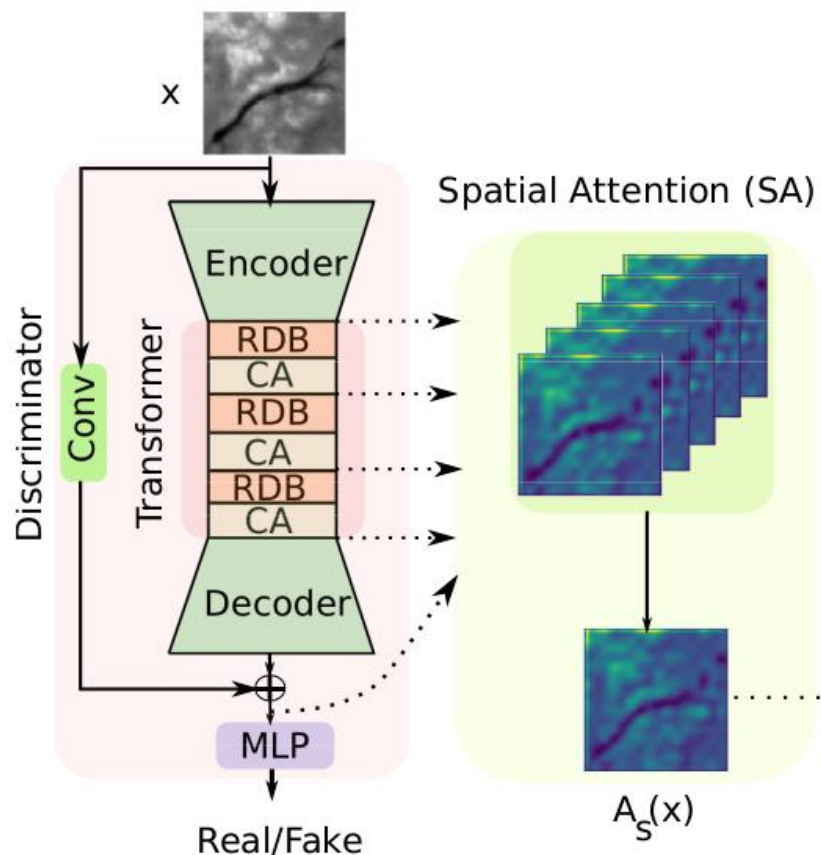
$$\mathcal{L}_{sa} = \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}, y \sim \mathbb{P}_y} \left[\|A_s(\hat{x}) - A_s(y)\|_2^2 \right]$$

Domain Adaptation Loss

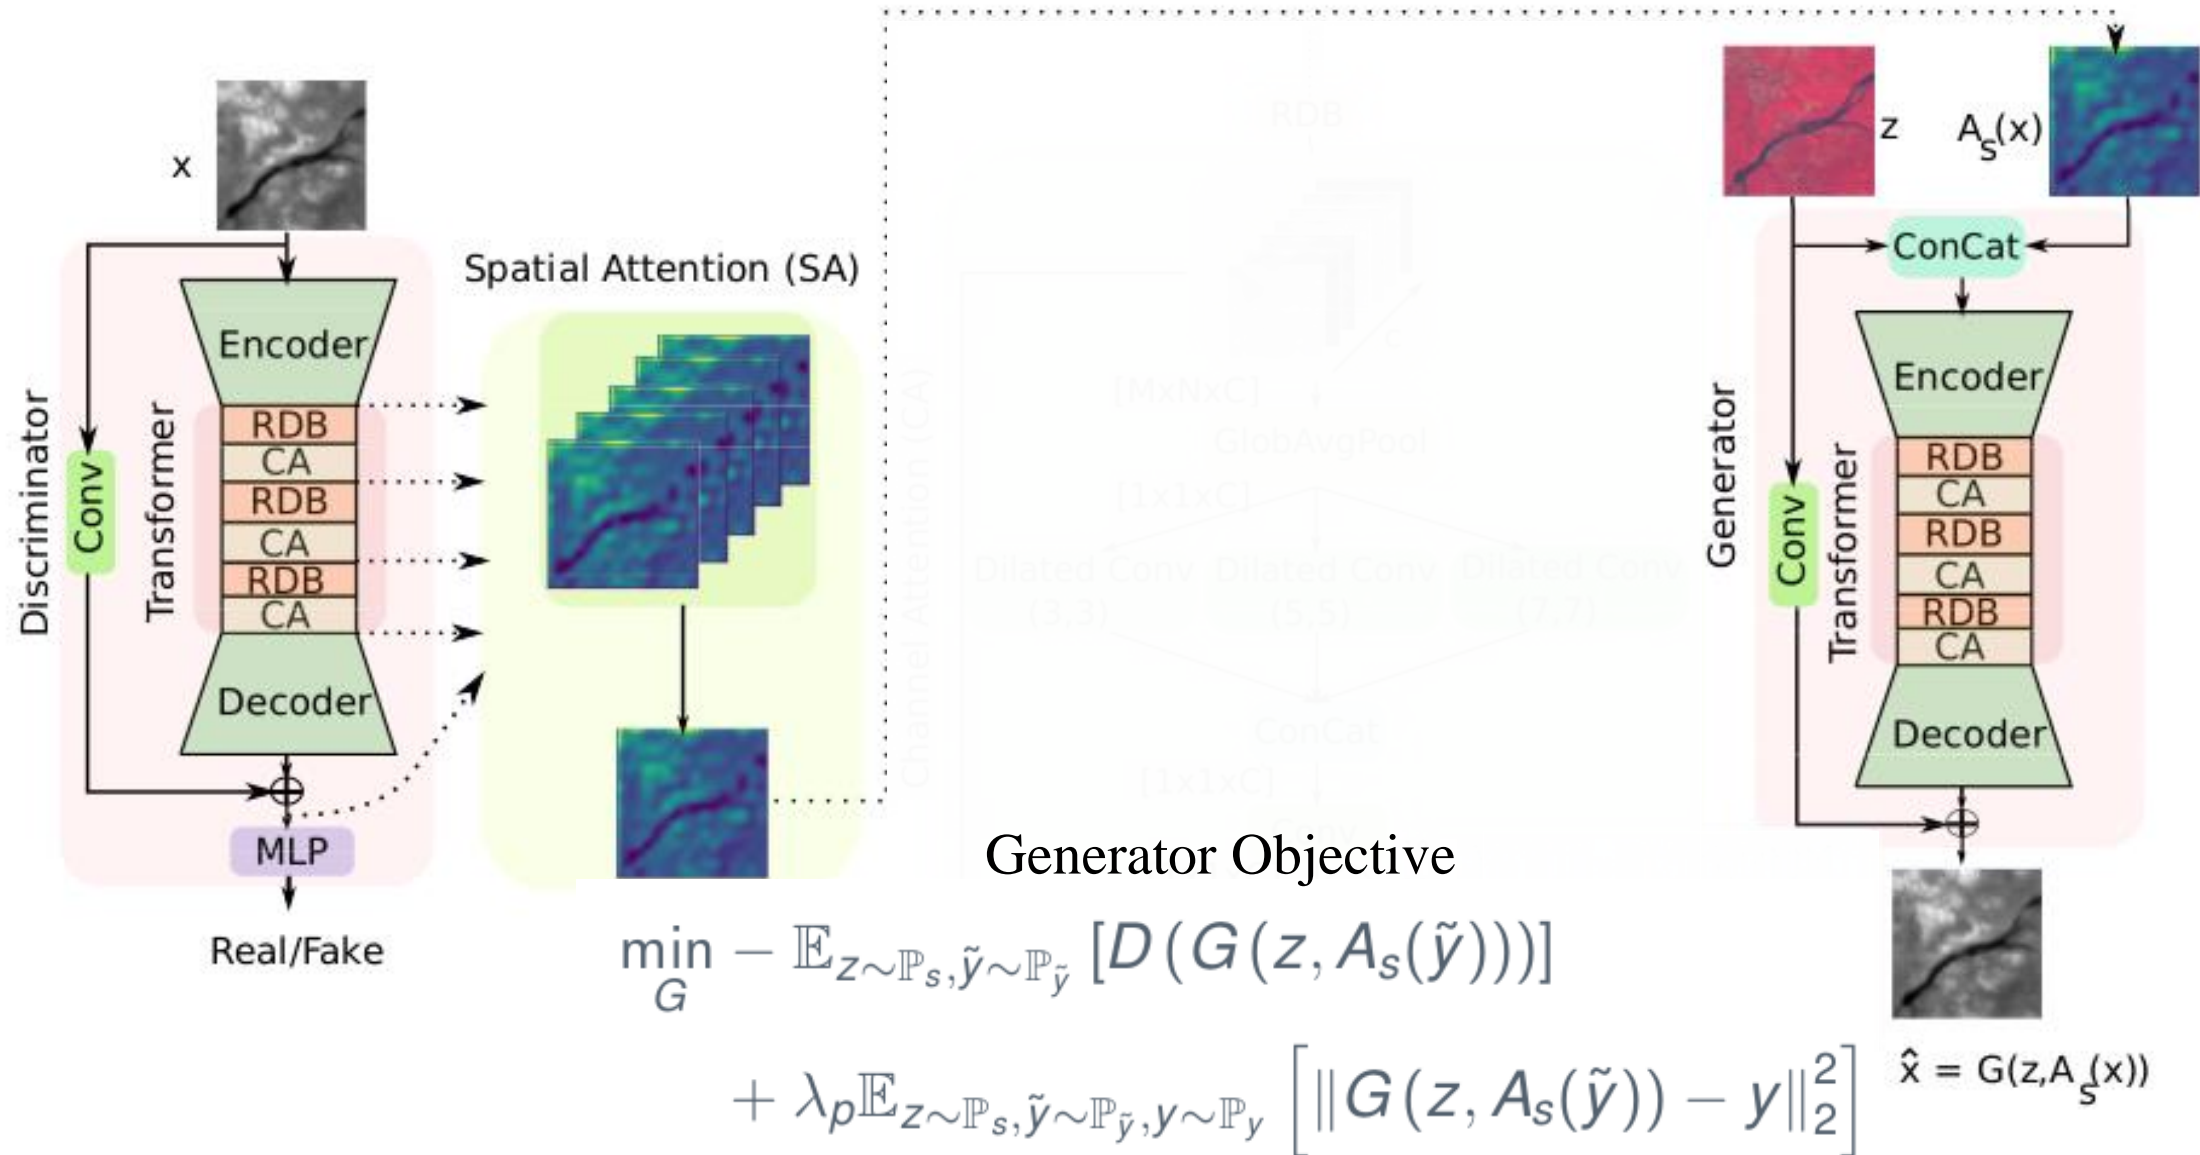
$$\mathcal{L}_{da} = \mathbb{E}_{\tilde{y} \sim \mathbb{P}_{\tilde{y}}, y \sim \mathbb{P}_y} \left[\|A_s(\tilde{y}) - A_s(y)\|_2^2 \right]$$

Discriminator Objective

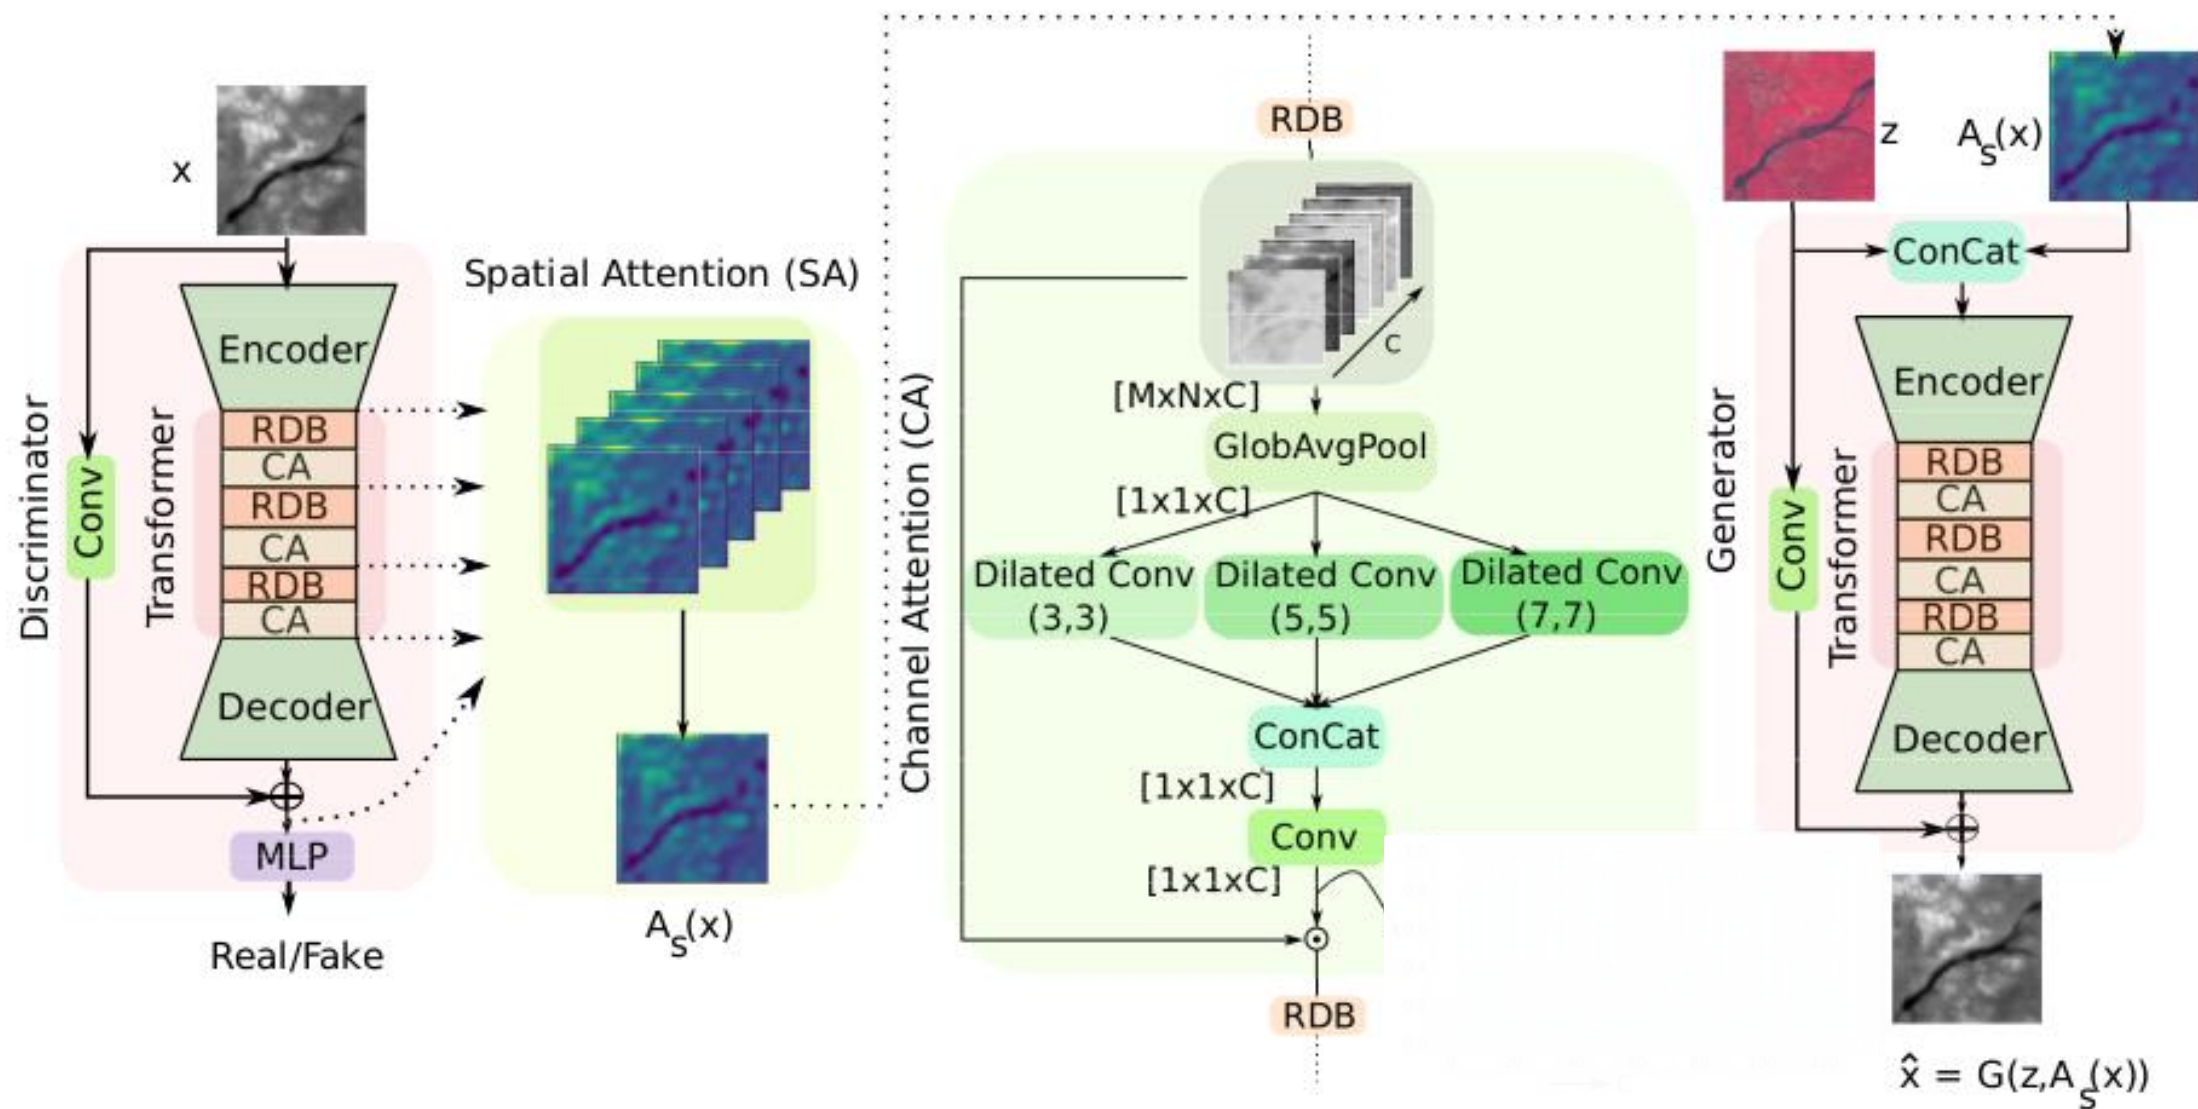
$$\begin{aligned} \min_D \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [D(\hat{x})] - \mathbb{E}_{x \sim \mathbb{P}_x} [D(x)] \\ + \lambda_{gp} \mathbb{E}_{\tilde{x} \sim \mathbb{P}_{\tilde{x}}} \left[(\|\nabla_{\tilde{x}} D(\tilde{x})\|_2 - 1)^2 \right] \\ + \lambda_{sa} \mathcal{L}_{sa} + \lambda_{da} \mathcal{L}_{da}, \end{aligned}$$



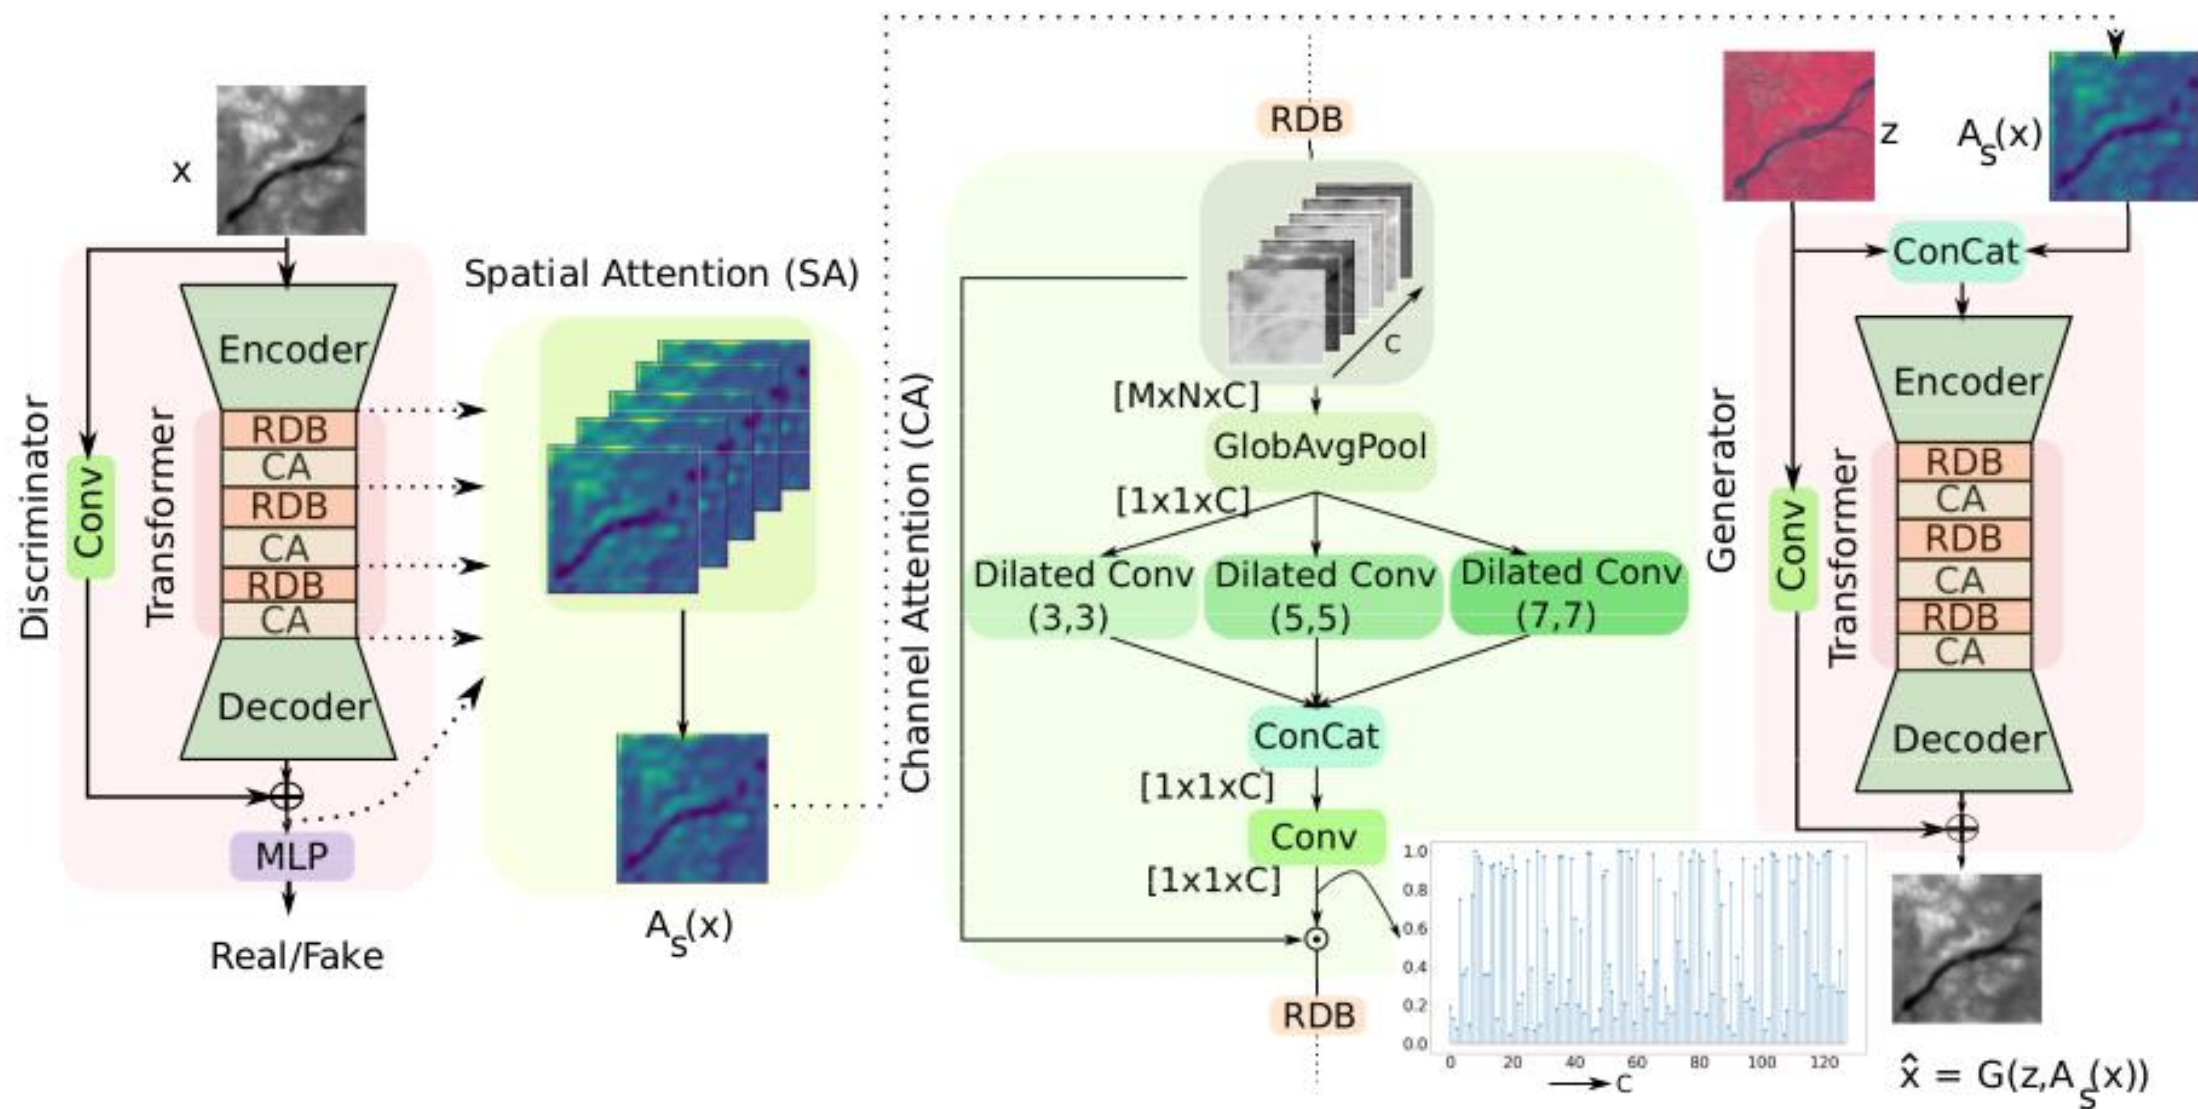
Spatial Attention from Discriminator



Spatio-Spectral Laplacian Attention




Spatio-Spectral Laplacian Attention




Spectral attention coefficients


Combining Spatial Attention with Source Bands




Multiplication:

- Attention module latches on to bright targets.
- Synthesized band contains blocky artifacts.






Method	RMSE	SSIM(%)	SRE(dB)	PSNR(dB)	SAM(deg)
AeroGAN [31]	21.62	86.03	44.62	36.50	12.15
DSen2 [21]	14.14	93.85	50.04	41.94	7.88
DeepSWIR [33]	13.75	94.02	50.35	42.27	7.66
ALERT [32]	12.97	94.54	50.81	42.80	7.48
S2A (ours)	11.74	95.08	50.83	42.76	6.87



- Learns to attend to relevant parts of source imagery.
- Homogeneous and heterogeneous targets are discernible.
- Similar features have similar attention coefficients

Wetland Delineation





— NIR

— Synthesized
SWIR

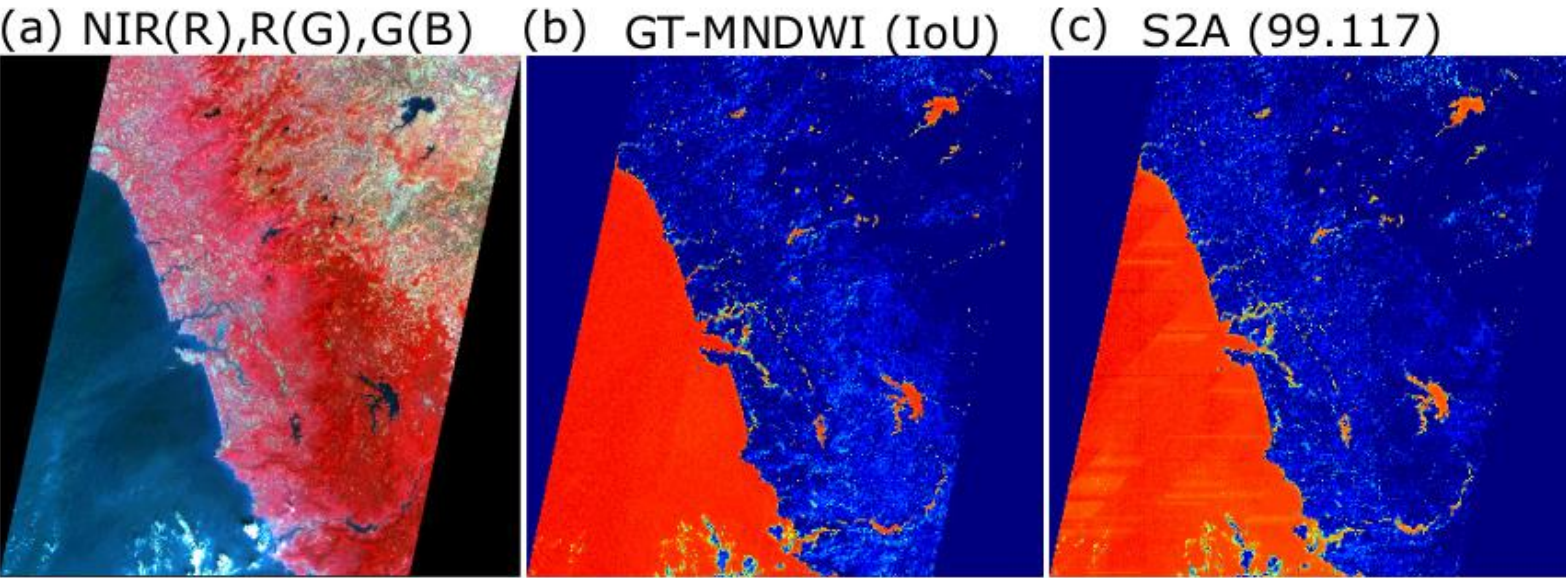
Water Segmentation




Wetland Delineation



Water Segmentation



Additional Value Product Generation




India

Hilly Terrain



Desert



Main land



Coastal

Overview

- Formulated super resolution as conditional band synthesis
- Regulated band synthesis through spatial and Laplacian spectral channel attention
- Introduced two new cost functions for the discriminator:
 - ◆ Spatial attention loss
 - ◆ Domain adaptation loss
- Experimented on multiple datasets:
 - ◆ LISS-3
 - ◆ LISS-4
 - ◆ WorldView-2
- Demonstrated real world applications of synthesized band:
 - ◆ Wetland delineation
 - ◆ Index based water segmentation
 - ◆ Additional value product generation/ Large area mosaic

Summary

- Three Pillars of Deep Learning
 - Setting Up DL Environment
 - Data Processing
 - Network Design
 - Visualization
 - Defining Problem Statement
 - Paired Training Data
 - Gradient Descent (GD)
 - Stochastic Gradient Descent (SGD)
 - Implementation Details
 - Linear Function Approximator
 - One Layer Neural Network Function Approximator
 - Two Layer Neural Network Function Approximator
 - Three Layer Convolutional Neural Network Function Approximator
 - Real World Application
 - Super Resolution
 - Multi-Spectral Band Synthesis